

Locks

Arnav Gupta

October 28, 2024

Contents

1	Lock Taxonomy	1
2	Spin Lock	2
2.1	Implementation	2
3	Blocking Locks	3
3.1	Mutex Lock	3
3.1.1	Implementation	3
3.1.2	uOwnerLock	6
3.1.3	Stream Locks	7
3.2	Synchronization Lock	7
3.2.1	Implementation	7
3.2.2	uCondLock	9
3.2.3	Programming Pattern	9
3.3	Barrier	9
3.3.1	Fetch Increment Barrier	10
3.3.2	uBarrier	10

1 Lock Taxonomy

Locks can be either spinning or blocking.

Spinning locks busy wait until an event occurs, so the task oscillates between ready and running due to time slicing (can also optionally yield).

Blocking locks do not busy wait but block until an event occurs, so some other mechanism must unblock the waiting task when the event happens.

2 Spin Lock

When a task is busy waiting, it usually loops until the critical section is unlocked or it is preempted and placed on the ready queue.

To increase uniprocessor efficiency, a task can:

- explicitly terminate its time slice
- move back to the ready state after only one event check fails

`yield` relinquishes the time-slice by rescheduling the running task back onto the ready queue. To increase multiprocessor efficiency, a task can yield after N event checks fail.

Adaptive Spin Lock: allow adjustment of spin duration

Most spin-lock implementations can cause starvation of 1+ tasks.

Spin lock is appropriate and necessary in situations where this is no other work to do.

2.1 Implementation

μ C++ provides `uSpinLock` (non-yielding) and `uLock` (yielding)

```
class uSpinLock {
public:
    uSpinLock(); // open
    void acquire();
    bool tryacquire(); // makes only 1 attempt to acquire lock
    void release();
}

class uLock {
public:
    uLock( unsigned int value = 1 );
    void acquire();
    bool tryacquire(); // makes only 1 attempt to acquire lock
    void release();
}
```

Locks built from atomic hardware instructions. Locks start closed (0) or opened (1).

Starvation can occur in theory, but rare in practice.

Cannot assign to or copy a lock, must use as reference/pointer. Can be used for synchronization or mutual exclusion.

3 Blocking Locks

For blocking locks, the acquiring task is not the only one responsible for detecting an open lock:

- acquiring task makes 1 check for open lock and then blocks
- releasing task is responsible for detecting a blocked acquirer and transferring or releasing the lock

Blocking locks reduce busy waiting through **cooperation**.

All blocking locks have

- state to facilitate lock semantics
- list of blocked acquirers

3.1 Mutex Lock

Used solely to provide mutual exclusion:

- separates lock usage between synchronization and mutual exclusion
- permits optimizations and checks as the lock only provides one specialized function

Mutex locks can be **single acquisition** (acquirer cannot acquire again) or **multiple acquisition** (acquirer can acquire multiple times). Multiple acquisition locks are called **owner locks**.

Owner locks can handle looping/recursion and may require 1 release or as many as it requires.

3.1.1 Implementation

```
class MutexLock {
    bool avail;
    Task * owner;
    queue<Task> blocked;
```

```

        SpinLock lock;
public:
    MutexLock() : avail { true }, owner{ nullptr } {}
    void acquire() {
        lock.acquire();
        while ( !avail && owner != currThread() ) {
            yieldNoSchedule( lock );
            lock.acquire();
        }
        avail = false;
        owner = currThread();
        lock.release();
    }

    void release() {
        lock.acquire();
        if ( owner != currThread() ) {
            // error check
        }
        owner = nullptr;
        if ( !blocked.empty() ) {
            // remove task from blocked list and make ready
        }
        avail = true;
        lock.release();
    }
}

```

`yieldNoSchedule` yields the processor time slice but does not reschedule thread to the ready queue.

This code allows for barging since `avail` and the lock reset.

For **barging avoidance**, must hold `avail` between releasing and unblocking task (protects against bounded overtaking):

```

class MutexLock {
    bool avail;
    Task * owner;
    queue<Task> blocked;
    SpinLock lock;
public:

```

```

MutexLock() : avail { true }, owner{ nullptr } {}
void acquire() {
    lock.acquire();
    if ( !avail && owner != currThread() ) {
        yieldNoSchedule( lock );
        // do not reacquire lock since avail == false
    } else {
        avail = false;
        lock.release();
    }
    owner = currThread();
}

void release() {
    lock.acquire();
    owner = nullptr;
    if ( !blocked.empty() ) {
        // remove task from blocked list and make ready
    } else {
        avail = true;
    }
    lock.release();
}
}

```

With these changes, bargers enter mutual exclusion protocol but block so the released task does not busy wait. This means mutual exclusion conceptually passed from releasing to unblocking tasks.

For **barging prevention**, must hold lock between releasing and unblocking task (protects against unbounded overtaking):

```

class MutexLock {
    bool avail;
    Task * owner;
    queue<Task> blocked;
    SpinLock lock;
public:
    MutexLock() : avail { true }, owner{ nullptr } {}
    void acquire() {
        lock.acquire();

```

```

        if ( !avail && owner != currThread() ) {
            yieldNoSchedule( lock );
            // do not reacquire lock
        } else {
            avail = false;
        }
        owner = currThread();
        lock.release();
    }

    void release() {
        lock.acquire();
        owner = nullptr;
        if ( !blocked.empty() ) {
            // remove task from blocked list and make ready
            // do not release lock
        } else {
            avail = true;
            lock.release();
        }
    }
}

```

The critical section is not bracketed by the spin lock when the lock is passed, so this works.

An alternative would be to leave the lock owner at the front of the blocked list to act as availability and owner variable. If the critical section is acquired, the blocked list must have a node on it to check if it is in use.

3.1.2 uOwnerLock

```

class uOwnerLock {
public:
    uOwnerLock();
    uBaseTask * owner(); // nullptr if no owner, else owner address
    unsigned int times(); // number of times lock has been acquired by owner
    void acquire();
    bool tryacquire();
    void release();
}

```

Must release as many times as acquired.

Can use `_Finally` or RAII to ensure lock is always released. However, this cannot be used for barging prevention.

3.1.3 Stream Locks

Can use `osacquire` for output streams and `isacquire` for input streams to ensure predictable streams.

3.2 Synchronization Lock

Used solely to block tasks waiting for synchronization.

Only state is a list of blocked tasks, so:

- acquiring task always blocks (no state to make it conditional)
- release is lost when no waiting task (no state to remember it)

Uses wait for acquire and signal for release.

3.2.1 Implementation

Needs mutual exclusion for safe implementation.

Location of mutual exclusion classifies synchronization lock:

- **external lock:** use an external lock to protect task list
- **internal lock:** use an internal lock to protect state

With external locking:

```
class SyncLock {
    Task * list;
public:
    SyncLock() : list{ nullptr } {}
    void acquire( MutexLock & m ) {
        // add self to task list
        yieldNoSchedule( m );
        // possibly reacquire mutex lock
    }
    void release() {
        if ( list != nullptr ) {
            // remove task from blocked list and make it ready
        }
    }
};
```

```

        }
    }
}

```

Uses external task to avoid lost release and need mutual exclusion to protect task list (and possible external state). The releasing task detects a blocked task and performs necessary cooperation.

To use with barging avoidance:

```

m.acquire();
if ( !s.empty() ) {
    s.release();
} else {
    occupied = false;
}
m.release();

```

To use with barging prevention:

```

m.acquire();
if ( !s.empty() ) {
    s.release();
} else {
    occupied = false;
    m.release();
}

```

With internal locking:

```

class SyncLock {
    Task * list;
    SpinLock lock;
public:
    SyncLock() : list{ nullptr } {}
    void acquire( MutexLock & m ) {
        lock.acquire();
        // add self to task list
        m.release();
        // can be interrupted here, but fine since spin lock
        yieldNoSchedule( lock );
        m.acquire(); // possibly reacquire mutex lock
    }
}

```



```

        void release() {
            lock.acquire();
            if ( list != nullptr ) {
                // remove task from blocked list and make it ready
            }
            lock.release();
        }
    }
}

```

This still takes an external lock for barging avoidance/prevention.

3.2.2 uCondLock

```

class uCondLock {
public:
    uCondLock();
    void wait( uOwnerLock & lock );
    bool signal(); // unblocks in FIFO order
    bool broadcast(); // unblocks all waiting tasks
    bool empty(); // are blocked tasks on the queue?
}

```

`wait` atomically blocks the calling task and releases the argument owner lock, and reacquires it before returning.

`signal` and `broadcast` do nothing for an empty conditional and return false, otherwise return true.

3.2.3 Programming Pattern

Must provide external mutual exclusion and protect against lost signal (release).

Should surround conditional lock with mutex lock.

3.3 Barrier

Coordinates a group of tasks performing a concurrent operation surrounded by sequential operations. Meant for synchronization, not mutual exclusion.

Two kinds of barriers:

- threads == group size

- threads > group size

Barrier retains state about the events it manages (num tasks blocked on the barrier). Most barriers use internal locking.

Barrier blocks each task at call to `block` until all tasks have called `block`. Last task to call `block` does not block, instead it releases all other tasks.

Must specify in advance number of `block` operations before tasks released.

Barriers are commonly used for synchronized one-shot and for synchronized start and end in a cycle.

Using a barrier is cheaper than creating and deleting tasks for each computation.

3.3.1 Fetch Increment Barrier

A spinning, $T == G$ barrier can be implemented with the fetch-increment instruction and a flag that waiters wait on.

3.3.2 uBarrier

μ C++ barrier is a blocking, $T > G$, barging-prevention coroutine, where `main` can be resumed by the last task arriving at the barrier:

```
#include <uBarrier.h>
_Cormonitor uBarrier {
    protected:
        void main() { for ( ;; ) suspend(); }
        virtual void last() { resume(); } // called by last task to barrier
    public:
        uBarrier( unsigned int total );
        unsigned int total() const; // # of tasks synchronizing
        unsigned int waiters() const; // # of waiting tasks
        void reset( unsigned int total ); // reset # of tasks synchronizing
        virtual void block(); //wait for Nth thread
}
```

`uBarrier` has implicit mutual exclusion so no barging.

Can build a barrier by inheriting from `uBarrier`, redefining `last`, `block`, and possibly `main`.

Coroutine barrier can be reused many times.